

Intro to JupyterHub on the National Research Platform

Presented by Cal Poly Humboldt - ITS

Brian Campbell - Academic Technology Specialist

Ravi Chalasani - ITS Research Support Technician



➤ National Research Platform

What is the National Research Platform?

- **Free** to access subject to accepting Acceptable Use Policy
- Collection of computing resources housed under one umbrella
- Funded by NSF
- Grew out of the former Pacific Research Platform
- Distributed network of over 400 servers with massive computing power - <https://dash.nrp-nautilus.io/>
- More computing power compared to running locally
- Instantly scalable to higher workloads



➤ **National Research Platform**

- **Do Not Mine crypto on NSF funded machines**
 - **federal crime**
- **Violators will get caught**
- **Violators will be reported to the FBI**
- **Cluster users may not use the computers to enrich themselves financially**
 - **Faculty/PIs: Please share this with your students**



➤ Accessing NRP Nautilus Portal

- URL for NRP Nautilus Portal: <https://nrp.ai/viz/resources/>
- We use the portal to create Namespaces for your projects
- We can add you and your students to your Namespace
- You can check for available resources
- NRP access is provided via Cyber Infrastructure CILogon and Authentik



➤ **NRP Support on Campus**

- Our process:
 1. Initial consultation about your project
 - Team based help with accessing NRP
 - Uploading any preexisting code/Notebooks to JupyterHub
 2. Troubleshooting, custom packages for specific needs
 - e.g. Otter-grader
 3. Uploading/downloading data to/from NRP
 4. Collaborate to adapt and improve processes
 5. Continued consultation until your project is successful



➤ How to Contact Us

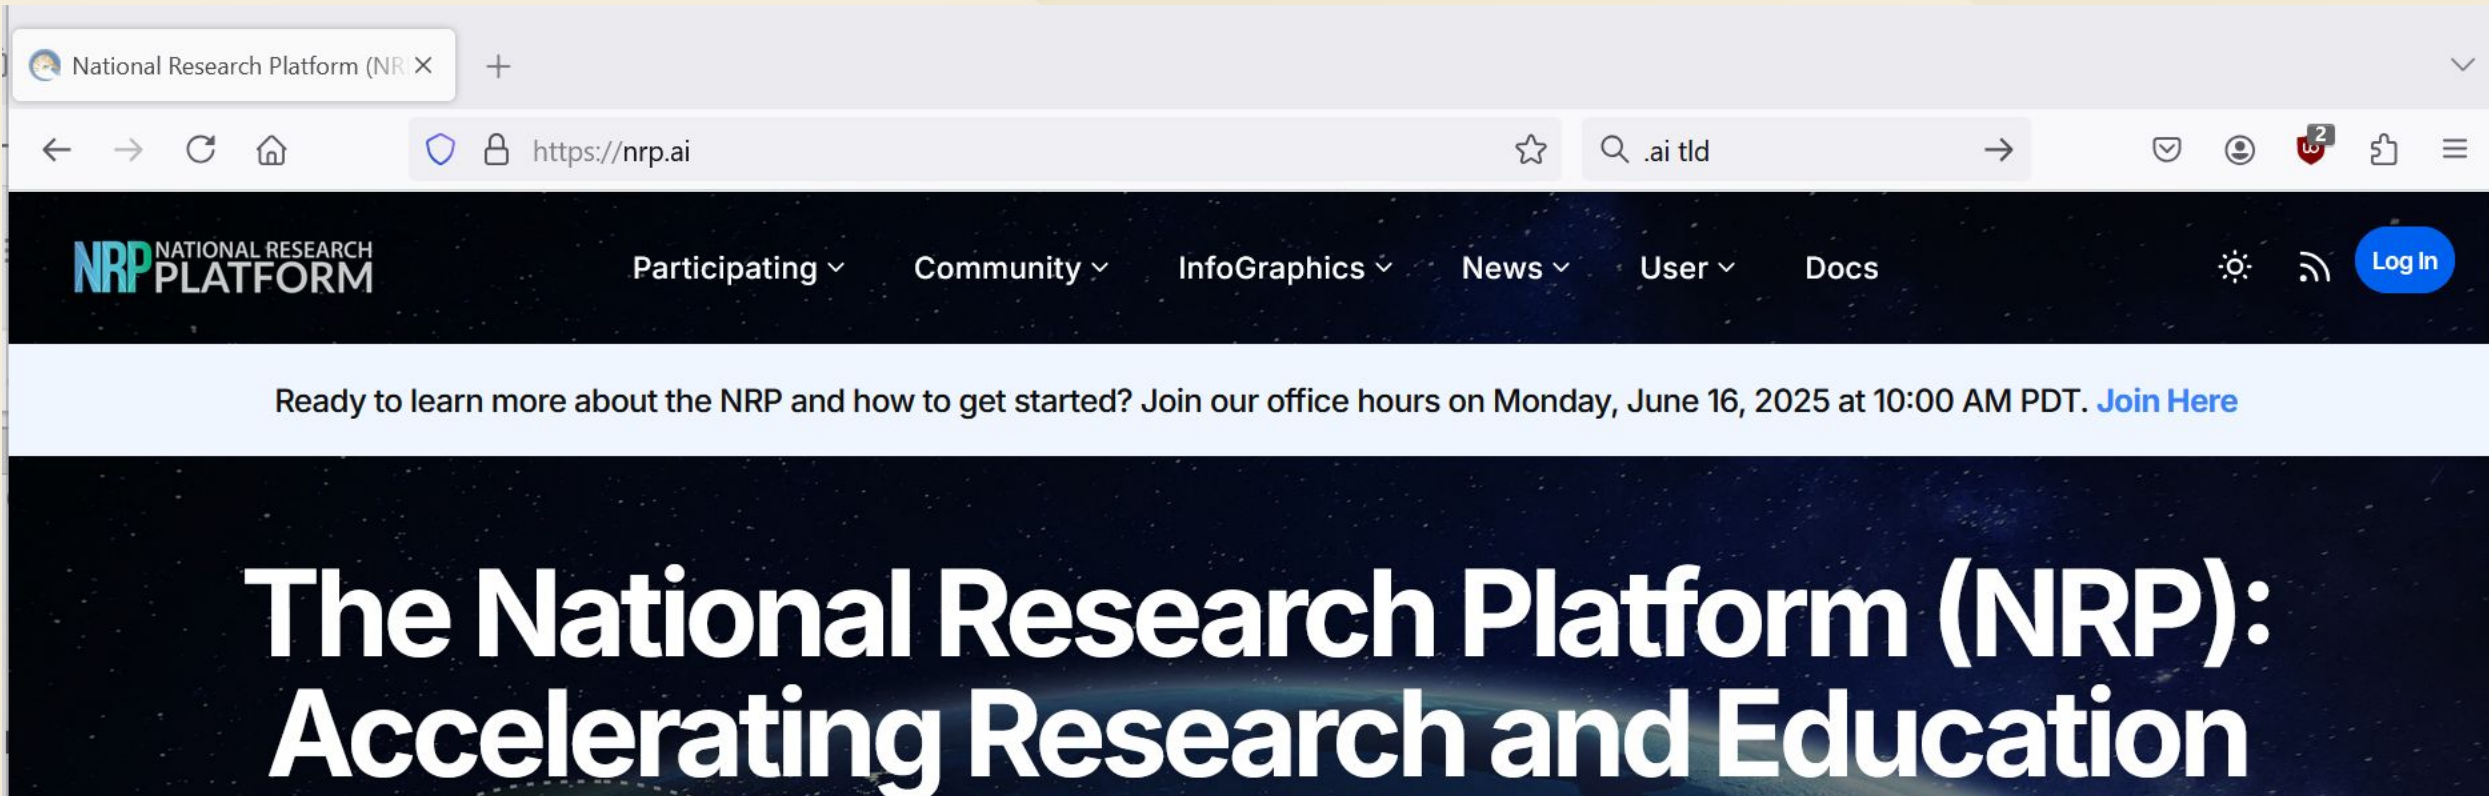
- SBS 413 (by appointment)
- Email: itsresearchsupport@humboldt.edu
- Create a HelpDesk ticket

☐ Research, Creative, and Scholarly Activities - Pre-Grant IT Dreaming, Planning, and/or Budgeting Consultation

- Website:
<https://its.humboldt.edu/research/high-performance-computing>
- Phone: x4100



> NRP Portal Pre-Login Screen



> NRP Portal CILogon Screen

Welcome To The CILogon Open X +

← → ↻ 🏠 🔒 https://cilogon.org/authorize?client_id=myproxy%3Aoa4mp%2C2012%3A%2Fcd ☆ 🔍 .ai tld

Nautilus requests access to the following information. If you do not approve this request, do not proceed.

- Your CILogon user identifier
- Your name
- Your email address
- Your username and affiliation from your identity provider

Select an Identity Provider

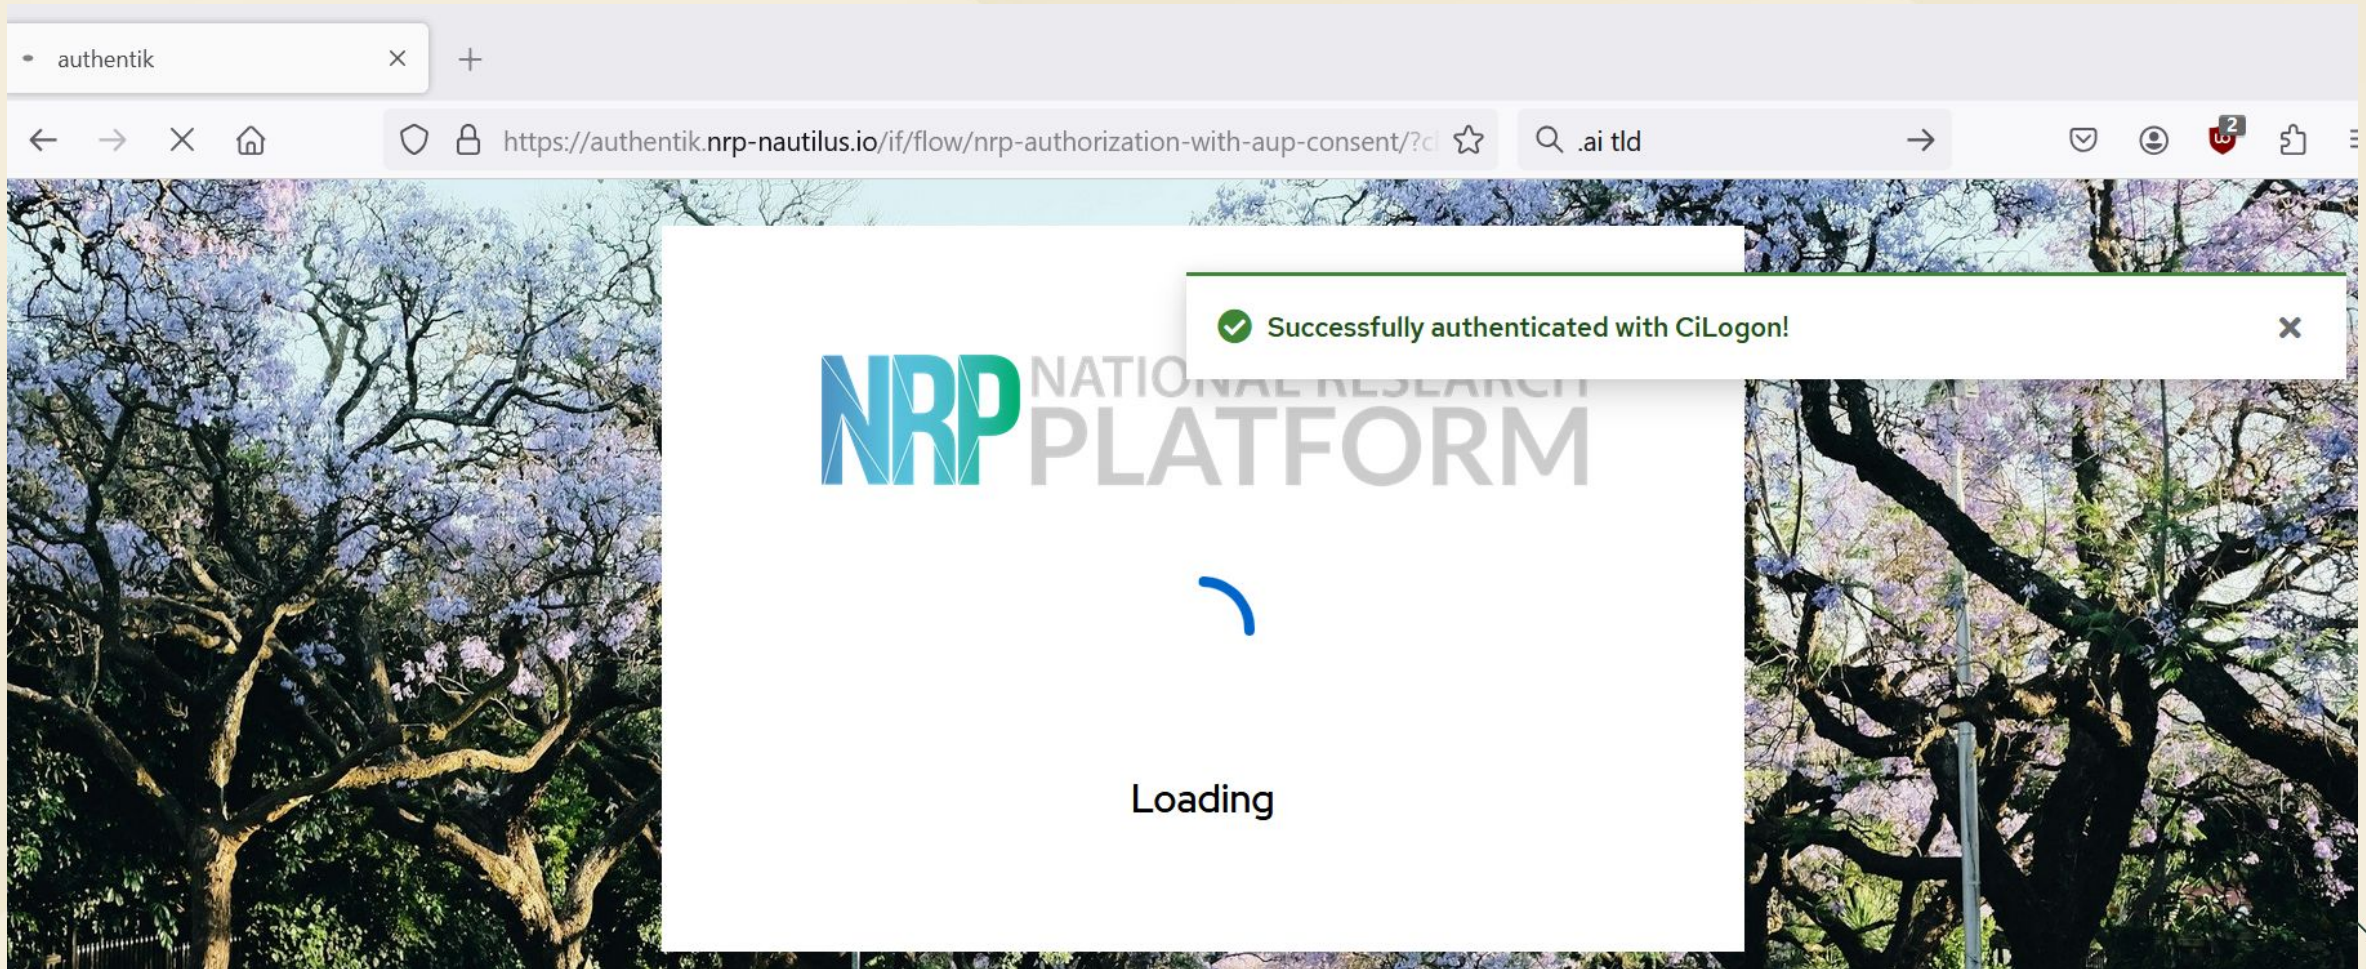
Cal Poly Humboldt ▲ ?

☐ Remember this selection ?

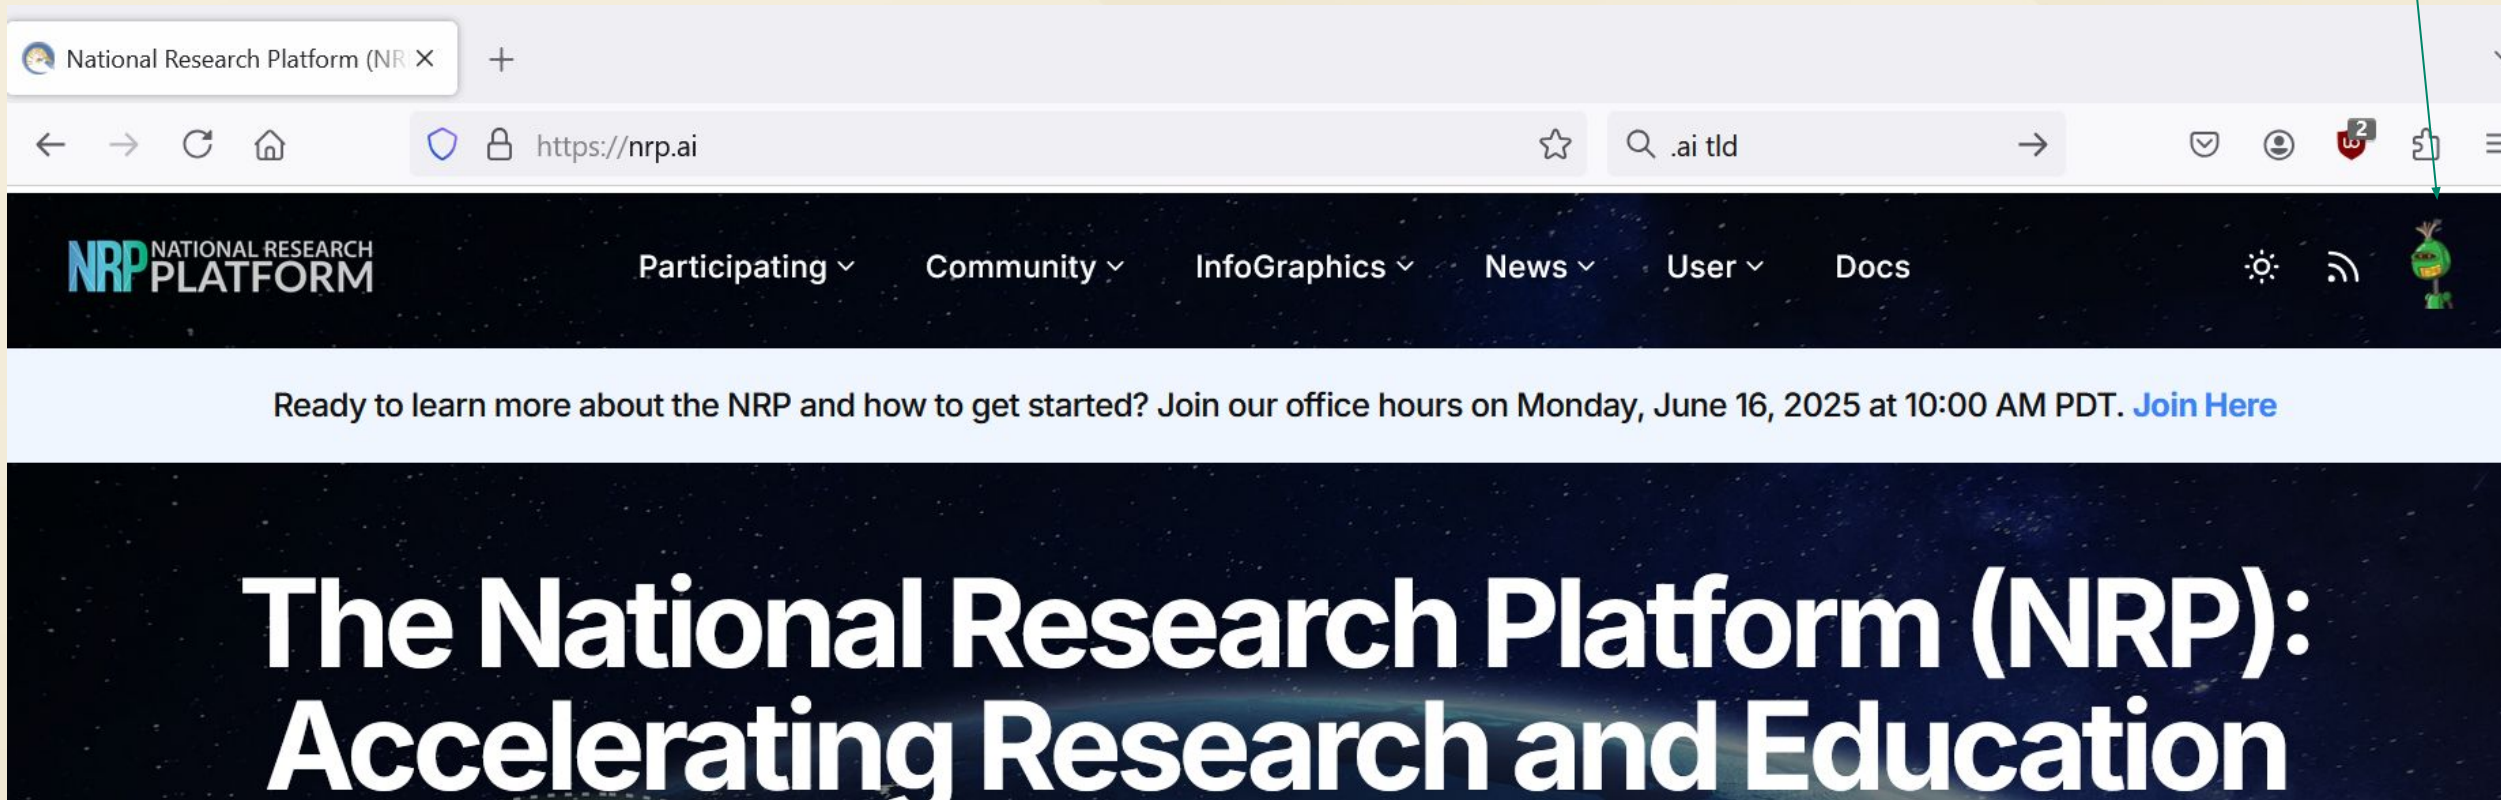
Log On



> NRP Portal Authentik Screen



> NRP Portal Logged-in Screen



The screenshot shows a web browser window with the address bar displaying "https://nrp.ai". The browser's tab bar shows "National Research Platform (NR X)". The website's header is dark blue with the "NRP NATIONAL RESEARCH PLATFORM" logo on the left. Navigation links include "Participating", "Community", "InfoGraphics", "News", "User", and "Docs". On the right of the header are icons for a sun, a Wi-Fi signal, and a small green alien character. Below the header is a light blue banner with the text: "Ready to learn more about the NRP and how to get started? Join our office hours on Monday, June 16, 2025 at 10:00 AM PDT. [Join Here](#)". The main content area has a dark blue background with the text: "The National Research Platform (NRP): Accelerating Research and Education".

National Research Platform (NR X)

https://nrp.ai

NRP NATIONAL RESEARCH PLATFORM

Participating Community InfoGraphics News User Docs

Ready to learn more about the NRP and how to get started? Join our office hours on Monday, June 16, 2025 at 10:00 AM PDT. [Join Here](#)

The National Research Platform (NRP): Accelerating Research and Education

➤ Jupyter... Hub, Notebook, etc

- JupyterLab: Local development environment for Jupyter Notebooks, allowing for markdown and code execution in “cells”
- It is not a programming language, it supports other languages
- JupyterHub (JH): Allows for execution of Jupyter Notebooks on a remote server
- Instantly transferable from local computer to high-performance computing resources
- Your data is stored in a 5GB Persistent Volume Claim (PVC)
- Code will keep running as long as your browser tab is open to JH
- Idle JH sessions will be terminated by NRP however your PVC data is preserved



> Local vs. NRP (with GPUs)

```
def batched_dot_mul_sum(a, b):  
    '''Computes batched dot by multiplying and summing'''  
    return a.mul(b).sum(-1)
```



2.879 s



0.000259 s

> How Do I Access JupyterHub?

- <https://jupyterhub-west.nrp-nautilus.io/>
- <https://github.com/cal-poly-humboldt/NRP-Tutorials>
- Use CyberInfrastructure Logon (CILogon)
 - Use Humboldt login



Select an Identity Provider

Cal Poly Humboldt ▲ ?

☐ Remember this selection ?

Log On

By selecting "Log On", you agree to the [privacy policy](#).



> JupyterHub Options

Server Options

By starting a jupyter instance you're agreeing to the [Acceptable Use Policy](#)

/home/jovyan is persistent volume, 5GB by default. Make sure you don't fill it up - jupyter won't start next time. You can ask admins to increase the size.

The storage is created in West ceph pool by default. You can ask admins to move it to a different region.

[Available resources page](#)

[GPU types guide](#)

Contact admins in [Matrix](#).

Region

GPUs

Cores

RAM, GB

GPU type

Refer to the [documentation](#) for images description.

Image
☐ Scipy
☐ R

- Select Region if desired
- GPU if needed
 - Graphics Processing Unit
 - additional brain for graphics
 - can be adapted for other compute
- Select CPUs called cores
 - Central Processing Unit
 - It's the brain of the computer
 - That's where the math is done



> JupyterHub Options

Server Options

By starting a jupyter instance you're agreeing to the [Acceptable Use Policy](#)

/home/jovyan is persistent volume, 5GB by default. Make sure you don't fill it up - jupyter won't start next time. You can ask admins to increase the size.

The storage is created in West ceph pool by default. You can ask admins to move it to a different region.

[Available resources page](#)

[GPU types guide](#)

Contact admins in [Matrix](#).

Region

GPUs

Cores

RAM, GB

GPU type

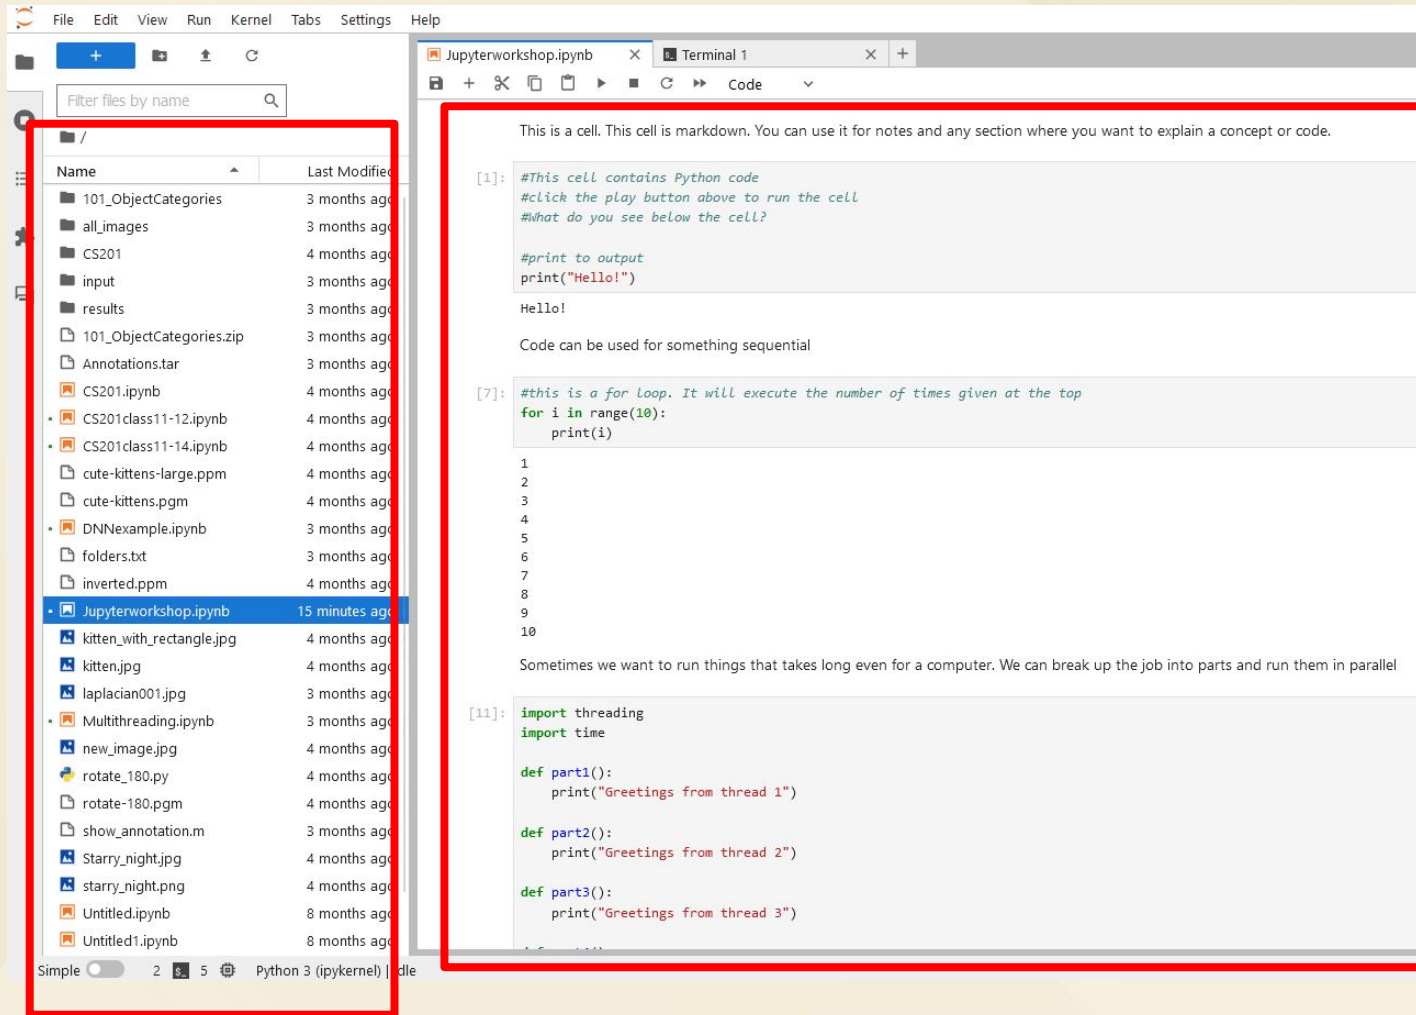
Refer to the [documentation](#) for images description.

Image
☒ Scipy
☐ R

- Select RAM
 - Random Access Memory
 - RAM larger than file size if possible
- Select Image (environment)
 - If don't know use default
 - Unless you are programming in R, then select R



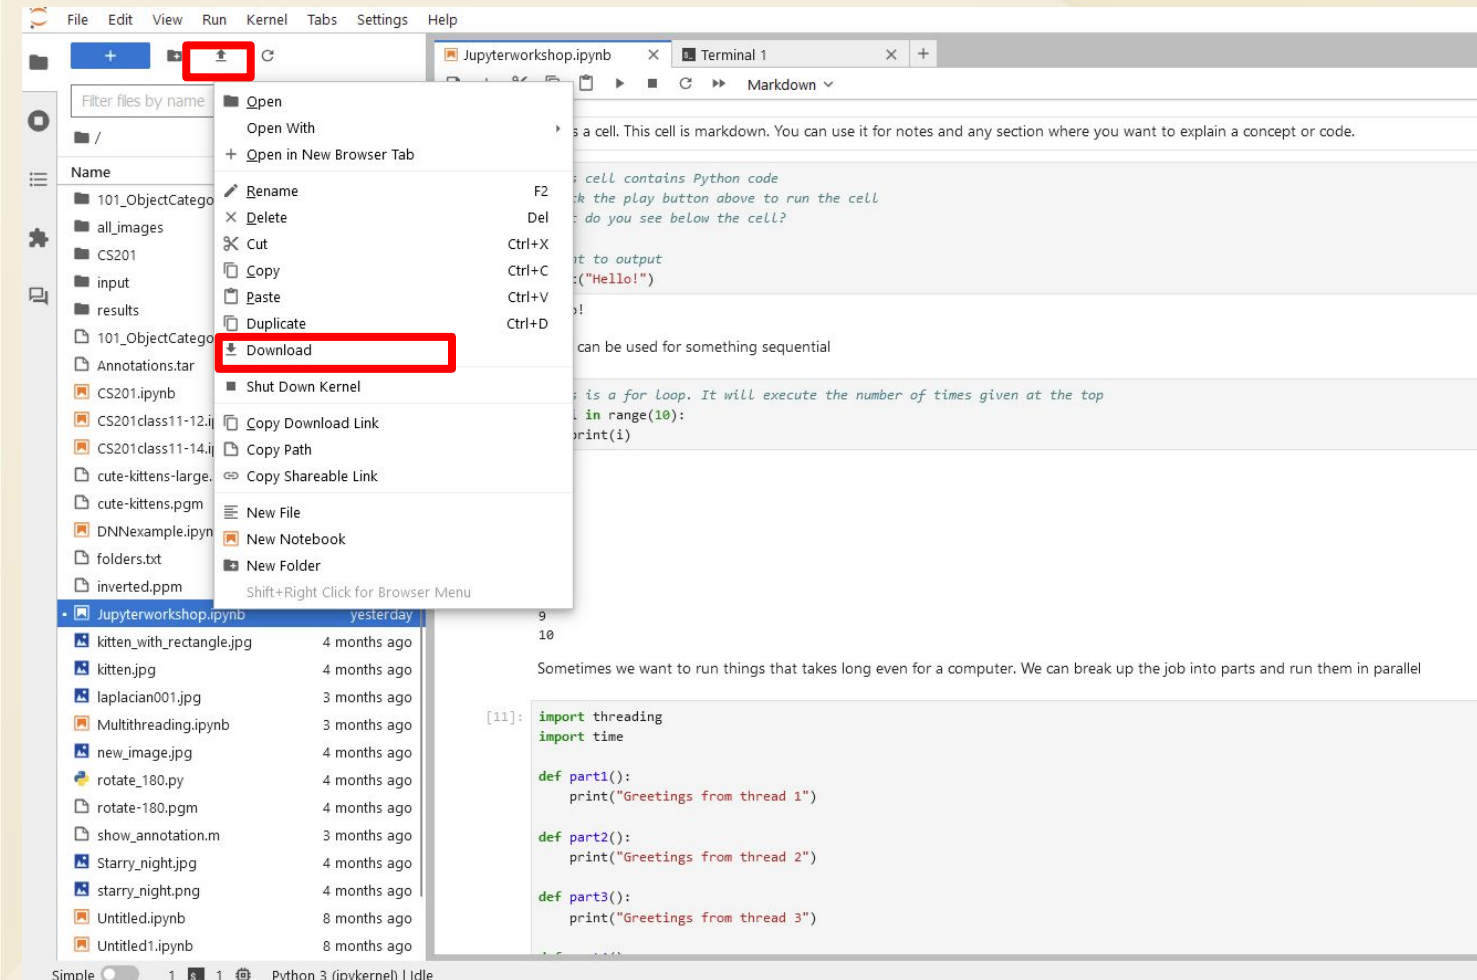
> JupyterHub Interface



- Graphical User Interface (GUI):
 - File system explorer on the left
 - Notebooks etc. on the right



JupyterHub Interface



- File System:
 - upload files
 - download files
 - R-click the file
 - select Download

JupyterHub Interface

The screenshot shows the JupyterHub interface with a file browser on the left and a notebook editor on the right. The notebook has three cells:

- Cell 1: A markdown cell containing the text: "This is a cell. This cell is markdown. You can use it for notes and any section where you want to explain a concept or code." This cell is highlighted with a red box.
- Cell 2: A code cell containing Python code:

```
[1]: #This cell contains Python code
#Click the play button above to run the cell
#What do you see below the cell?

#print to output
print("Hello!")

Hello!

Code can be used for something sequential
```

 This cell is highlighted with a red box.
- Cell 3: A code cell containing Python code:

```
[7]: #this is a for loop. It will execute the number of times given at the top
for i in range(10):
    print(i)

1
2
3
4
5
6
7
8
9
10
```

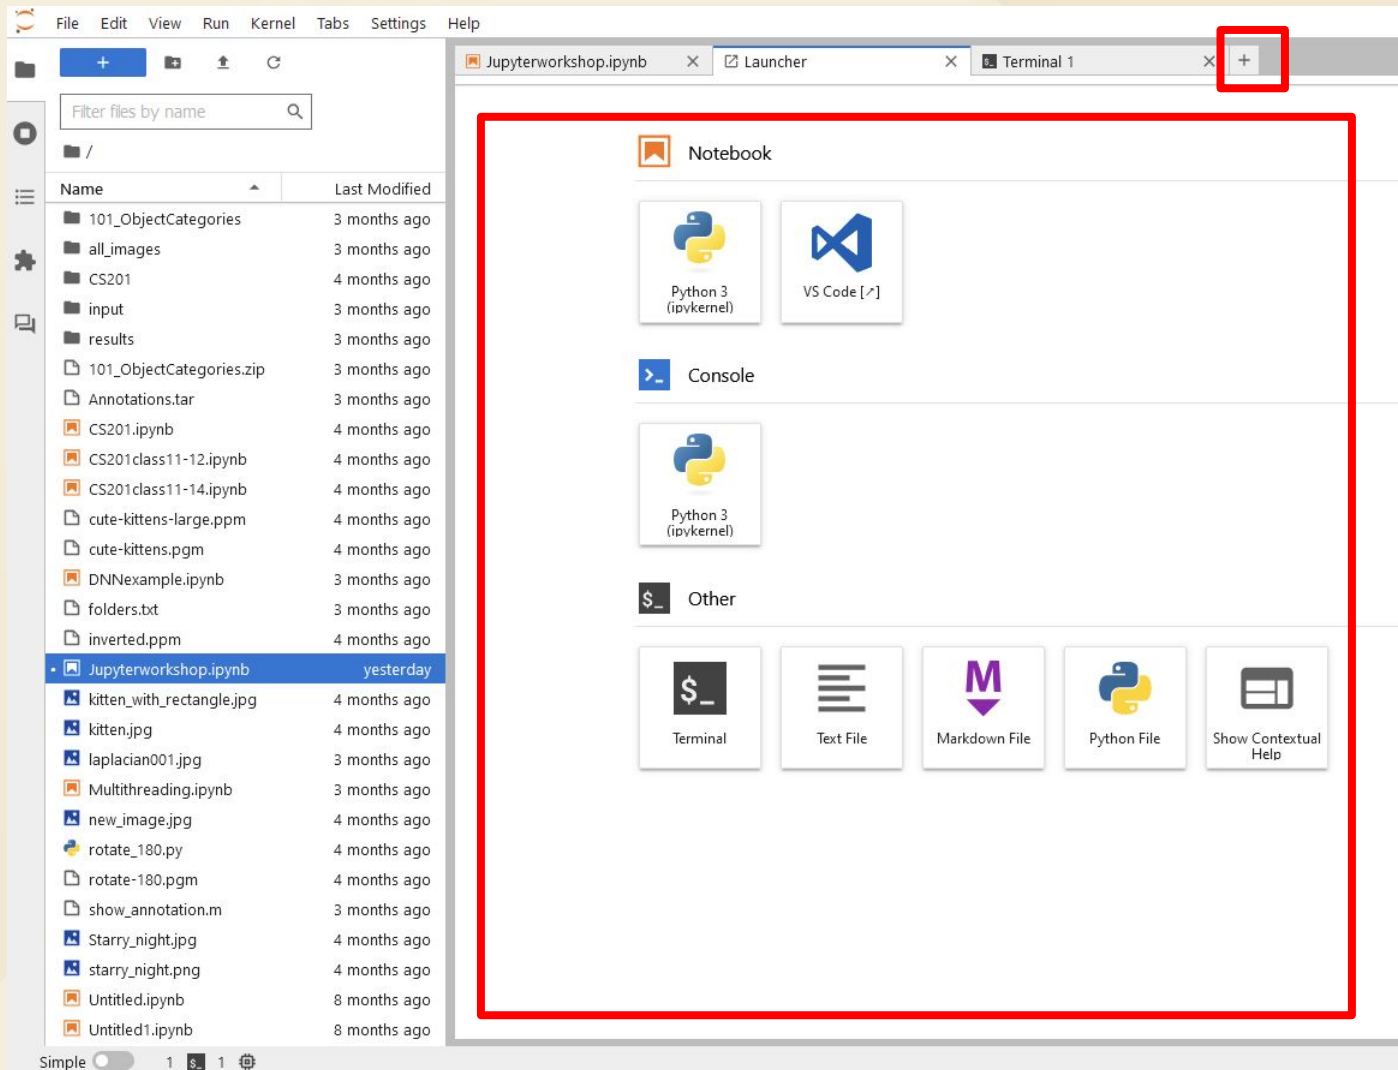
 This cell is highlighted with a red box. Below the code, the output area is highlighted with a red box, showing the numbers 1 through 10.

Red arrows point from the list on the right to these highlighted areas:

- Red arrow from "Notebook:" to the first cell.
- Red arrow from "Cells" to the second cell.
- Red arrow from "Markdown for notes" to the first cell.
- Red arrow from "Code for Python or R" to the second cell.
- Red arrow from "Output areas after the code runs" to the output area of the third cell.
- Red arrow from "Pick which cells to run" to the third cell.

- Notebook:
 - Cells
 - Markdown for notes
 - Code for Python or R
 - Output areas after the code runs
 - Pick which cells to run


JupyterHub Interface



- Click the + to use the Launcher
 - Open a New Tab
 - New Notebook File
 - Terminal Windows



JupyterHub Interface

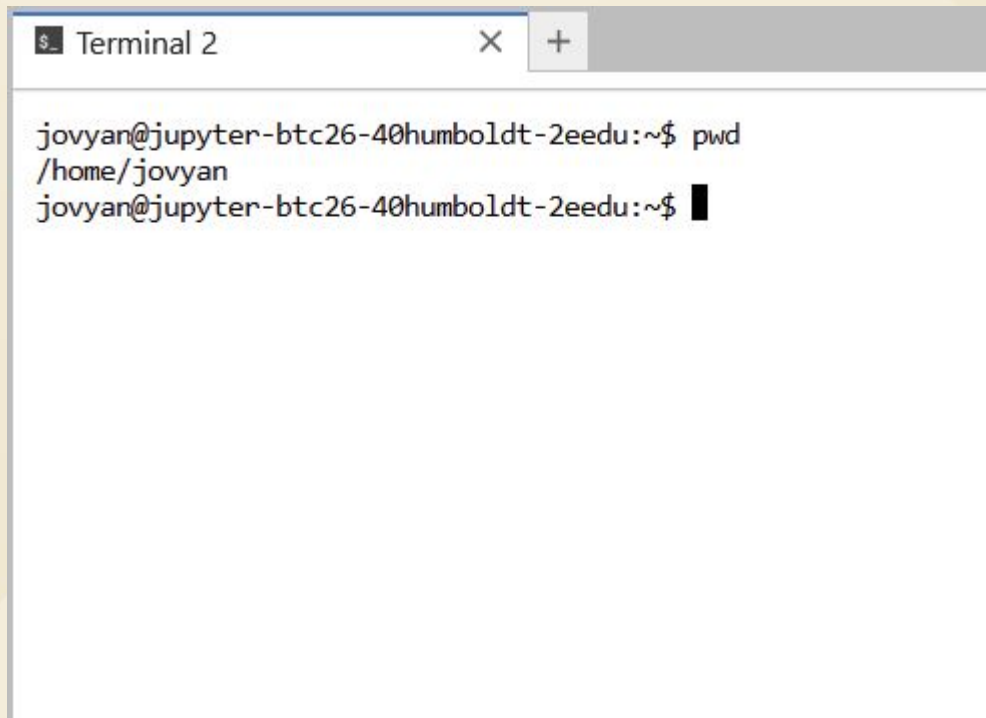


- Terminal:
 - Linux Command Line Interface
 - Usually optional
 - can use to install additional packages
 - Allows for installations
 - Command to try:
 - pwd

Interactive demo of JH



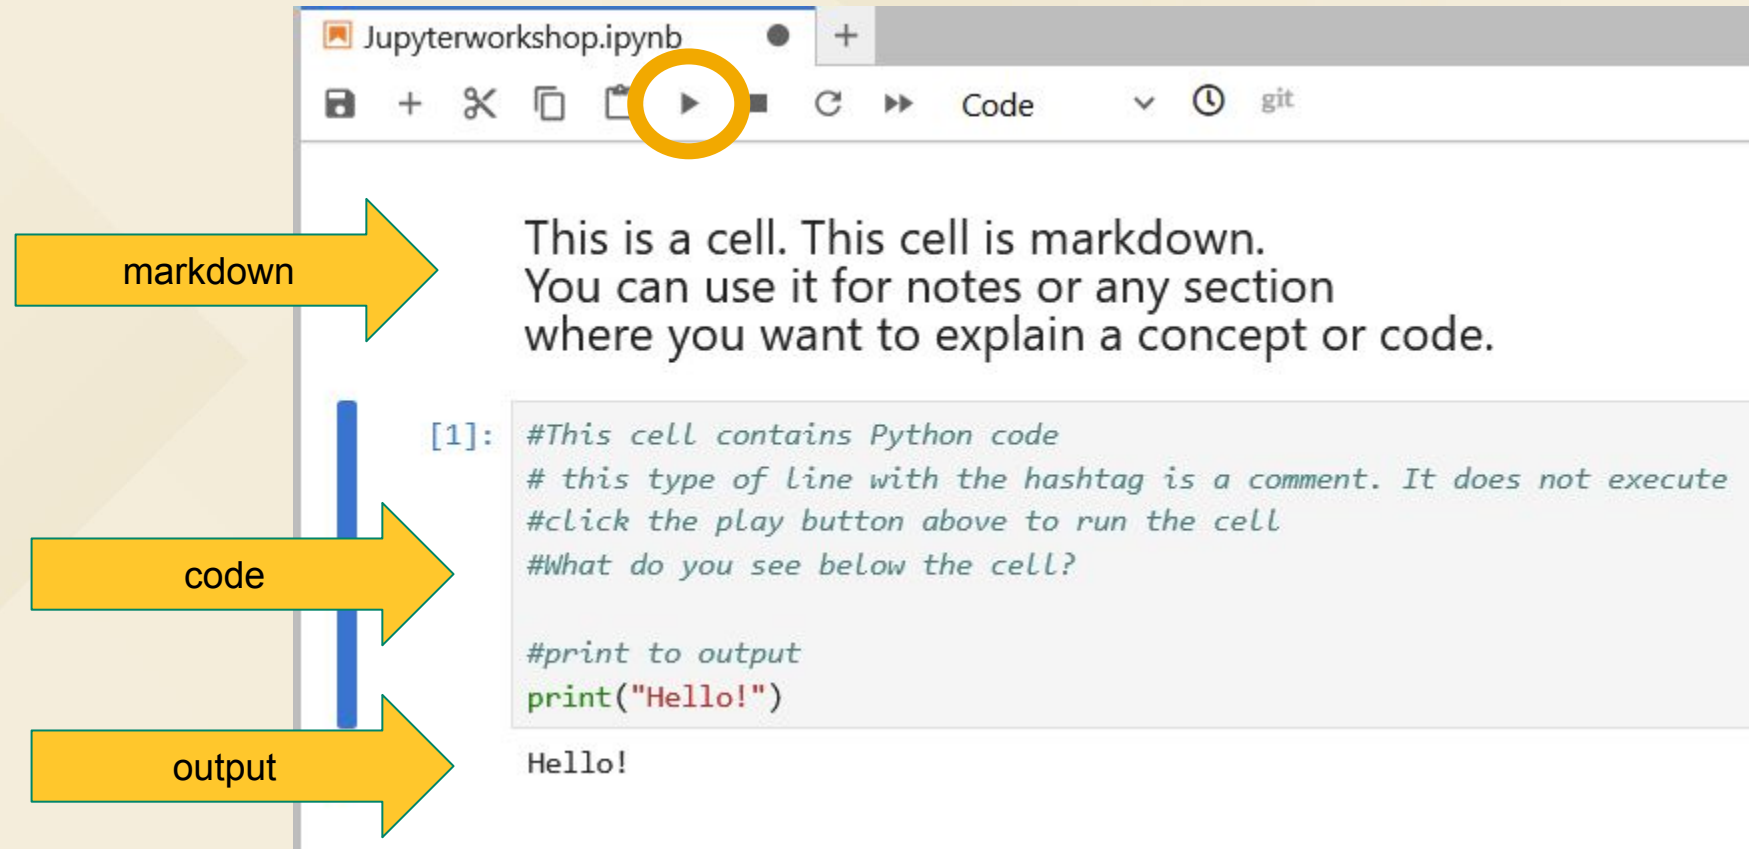
➤ JupyterHub Environment: Linux Terminal



```
Terminal 2
jovyan@jupyter-btc26-40humboldt-2edu:~$ pwd
/home/jovyan
jovyan@jupyter-btc26-40humboldt-2edu:~$
```

- Full Linux terminal
 - run commands
 - ex: pwd

➤ JupyterHub Environment: Notebook



The screenshot shows a Jupyter Notebook window titled "Jupyterworkshop.ipynb". The toolbar at the top includes icons for saving, adding, deleting, copying, pasting, and running code. The run button (a play icon) is circled in orange. Below the toolbar, there are three yellow arrows pointing to different parts of the notebook:

- An arrow labeled "markdown" points to a text cell containing: "This is a cell. This cell is markdown. You can use it for notes or any section where you want to explain a concept or code."
- An arrow labeled "code" points to a code cell containing:

```
[1]: #This cell contains Python code  
# this type of line with the hashtag is a comment. It does not execute  
#click the play button above to run the cell  
#What do you see below the cell?  
  
#print to output  
print("Hello!")
```
- An arrow labeled "output" points to the output area below the code cell, which displays "Hello!".



> Python Basics: Math



▼ You can do mathematical functions.

```
[3]: #You can do math in Python  
#What do you see below the cell?  
  
#add 1 + 1 and store the value in a variable called x  
x = 1 + 1  
  
#print to output  
print("the value of x is:", x)  
  
the value of x is: 2
```

› Python Basics: Looping

Code can be used for something sequential

```
[7]: #this is a for loop. It will execute the number of times given at the top  
for i in range(10):  
    print("Hello")  
    print(i) # must indent so Python knows it is part of the loop
```

#what is shown in the output?

```
Hello  
0  
Hello  
1  
Hello  
2  
Hello  
3  
Hello  
4  
Hello  
5  
Hello  
6  
Hello  
7  
Hello  
8  
Hello  
9
```



➤ Python Basics: Functions

You can call another piece of code called a function

```
[15]: #this part only calls when the code calls it  
def function1(x):  
    # must indent so Python knows it is part of the function  
    print("Greetings from function", x)  
  
#this is the main part of the code  
function1(1)
```

Greetings from function 1



> Python Basics: Libraries

In Python you can easily import libraries for additional functions

```
[17]: import threading  
import time
```



➤ Stress the CPU: Some Math in Serial

```
import time
import math

# a function with some math functions in a loop
def large_loop(limit):
    a = 1
    #note the nested indentations of both the function and the loop
    for i in range(limit):
        a = math.factorial(16)
        a = math.cbrt(139045)
        a = math.cos(float(i))
        a = math.cos(46.893)
        a = math.gamma(16.2)

#will run for 2^28 times
loop_limit = 2**28

#start the clock
start_time = time.time()
print("start time: ", start_time)

#call the function with the loop
large_loop(loop_limit) #pass the loop limit to the function

#stop the clock and print the time
finish_time = time.time() - start_time
print("The time it took to run is", finish_time, "seconds")
```



➤ Stress the CPU: Some Math in Serial

- Running 1 CPU only

```
import time
import math

# a function with some math functions in a loop
def large_loop(limit):
    a = 1
    #note the nested indentations of both the function and the loop
    for i in range(limit):
        a = math.factorial(16)
        a = math.cbrt(139045)
        a = math.cos(float(i))
        a = math.cos(46.893)
        a = math.gamma(16.2)

#will run for 2^28 times
loop_limit = 2**28

#start the clock
start_time = time.time()
print("start time: ", start_time)

#call the function with the loop
large_loop(loop_limit) #pass the loop limit to the function

#stop the clock and print the time
finish_time = time.time() - start_time
print("The time it took to run is", finish_time, "seconds")
```

```
top - 18:34:59 up 14 days, 18:10,  0 user,  load average: 11.19, 12.28, 11.86
Tasks:  7 total,   3 running,  4 sleeping,   0 stopped,   0 zombie
%Cpu(s): 17.2 us,  0.5 sy,  0.0 ni, 82.3 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
MiB Mem : 257797.2 total, 140105.8 free,  37235.6 used,  82726.7 buff/cache
MiB Swap:   0.0 total,    0.0 free,    0.0 used. 220561.6 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	%CPU	%MEM	TIME+	COMMAND
130	jovyan	20	0	1064212	77940	18896	100.0	0.0	3:16.59	python
30	jovyan	20	0	2150996	339824	75152	6.7	0.1	1:05.84	jupyterhub-sing
225	jovyan	20	0	11968	5420	3248	R 0.3	0.0	0:00.45	top
1	root	20	0	2692	1044	952	S 0.0	0.0	0:00.04	tini
7	root	20	0	14280	5740	4932	S 0.0	0.0	0:00.01	sudo
109	jovyan	20	0	760348	67932	18784	S 0.0	0.0	0:01.10	python
222	jovyan	20	0	7608	4264	3712	S 0.0	0.0	0:00.03	bash

➤ Stress the CPU: Some Math in Serial

- 198 seconds for 1 CPU



```
import time
import math

# a function with some math functions in a loop
def large_loop(limit):
    a = 1
    #note the nested indentations of both the function and the loop
    for i in range(limit):
        a = math.factorial(16)
        a = math.cbrt(139045)
        a = math.cos(float(i))
        a = math.cos(46.893)
        a = math.gamma(16.2)

#will run for 2^28 times
loop_limit = 2**28

#start the clock
start_time = time.time()
print("start time: ", start_time)

#call the function with the loop
large_loop(loop_limit) #pass the loop limit to the function

#stop the clock and print the time
finish_time = time.time() - start_time
print("The time it took to run is", finish_time, "seconds")
```

```
start time: 1744911558.9427457
The time it took to run is 198.4164820098877 seconds
```

> Multiprocessing

```
import multiprocessing
import time
import math

#functions with different math functions to stress the CPU(s)
def large_loop(stop, process):
    print("Greetings from process" , process)
    a = 1
    for i in range(stop):
        a = math.factorial(16)
        a = math.cbrt(139045)
        a = math.cos(float(i))
        a = math.cos(46.893)
        a = math.gamma(16.2)

# set the limit of the loop
limit = int((2**28)/4)

#print a greeting and start the clock
print("Greetings from the main process.")
start_time = time.time()
```

```
#set up all 4 processes
process1 = multiprocessing.Process(target = large_loop, args = (limit, 1))
process2 = multiprocessing.Process(target = large_loop, args = (limit, 2))
process3 = multiprocessing.Process(target = large_loop, args = (limit, 3))
process4 = multiprocessing.Process(target = large_loop, args = (limit, 4))

#starting then ending the processes
process1.start()
process2.start()
process3.start()
process4.start()
process1.join()
process2.join()
process3.join()
process4.join()

#print a greeting, stop the clock and print the time
print("It's the main process again!")
print("The time it took to execute is", time.time() - start_time, "seconds")
print("Processes 1-4 have finished executing.")
```

4 processes



start

stop

- break up the code into 4 parts and run in parallel

› Multiprocessing

- Running 4 CPUs simultaneously

```
top - 18:21:10 up 14 days, 17:56, 0 user, load average: 11.83, 11.35, 11.19
Tasks: 11 total, 5 running, 6 sleeping, 0 stopped, 0 zombie
%Cpu(s): 15.9 us, 1.1 sy, 0.0 ni, 83.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 257797.2 total, 141987.1 free, 35462.0 used, 82619.5 buff/cache
MiB Swap: 0.0 total, 0.0 free, 0.0 used, 222335.2 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
183	jovyan	20	0	711912	65464	65464	R	100.0	0.0	0:40.01	python
204	jovyan	20	0	711912	65516	65516	R	100.0	0.0	0:39.73	python
195	jovyan	20	0	711912	65472	65472	R	99.3	0.0	0:39.55	python
185	jovyan	20	0	711912	65464	65464	R	97.3	0.0	0:39.27	python
30	jovyan	20	0	2015120	334404	75152	S	0.0	0.1	0:25.64	jupyterhub-sing
1	root	20	0	2692	1044	952	S	0.0	0.0	0:00.02	tini
7	root	20	0	14280	5740	4932	S	0.0	0.0	0:00.01	sudo
109	jovyan	20	0	760348	67932	18784	S	0.0	0.0	0:01.02	python
130	jovyan	20	0	769284	77620	18832	S	0.0	0.0	0:01.21	python
222	jovyan	20	0	7608	4264	3712	S	0.0	0.0	0:00.03	bash
225	jovyan	20	0	11968	5420	3248	R	0.0	0.0	0:00.00	top



> Multiprocessing

```
import multiprocessing
import time
import math

#functions with different math functions to stress the CPU(s)
def large_loop(stop, process):
    print("Greetings from process" , process)
    a = 1
    for i in range(stop):
        a = math.factorial(16)
        a = math.cbrt(139045)
        a = math.cos(float(i))
        a = math.cos(46.893)
        a = math.gamma(16.2)

# set the limit of the loop
limit = int((2**28)/4)

#print a greeting and start the clock
print("Greetings from the main process.")
start_time = time.time()
```

```
#set up all 4 processes
process1 = multiprocessing.Process(target = large_loop, args = (limit, 1))
process2 = multiprocessing.Process(target = large_loop, args = (limit, 2))
process3 = multiprocessing.Process(target = large_loop, args = (limit, 3))
process4 = multiprocessing.Process(target = large_loop, args = (limit, 4))

#starting then ending the processes
process1.start()
process2.start()
process3.start()
process4.start()
process1.join()
process2.join()
process3.join()
process4.join()

#print a greeting, stop the clock and print the time
print("It's the main process again!")
print("The time it took to execute is", time.time() - start_time, "seconds")
print("Processes 1-4 have finished executing.")
```

```
Greetings from the main process.
Greetings from processGreetings from process 1
2
Greetings from processGreetings from process 34
```

```
It's the main process again!
The time it took to execute is 53.09485905075073 seconds
Processes 1-4 have finished executing.
```



- About ¼ the time

➤ **NRP Beyond JupyterHub**

Margarete Walden Fisheries Simulation

- Laptop - runtime is hours
- Jupyter - runtime a few minutes
- Kubernetes - increase parameters, automation, multiple parallel runs
- Python automated Kubernetes - further increased automation, expanded scope of research



➤ How to Contact Us

- SBS 413 (by appointment)
- Email: itsresearchsupport@humboldt.edu
- Create a HelpDesk ticket

☐ Research, Creative, and Scholarly Activities - Pre-Grant IT Dreaming, Planning, and/or Budgeting Consultation

- Website:
<https://its.humboldt.edu/research/high-performance-computing>
- Phone: x4100

